

Why the Commerce Protocol Draws Its Boundaries Where It Does

Version 1.0, 6 September 2026. A whitepaper for engineers who verify purchases on a server: why the OpenIAP Commerce Protocol draws its decision boundaries where it does, and what each one costs you to get wrong. It reports no measured result and has not been peer reviewed. [SPEC.md](#) is authoritative for the normative wording; this explains the reasoning behind it.

Published as a PDF at <https://www.openiap.dev/commerce-protocol-rationale.pdf>. Released under the MIT License, like the rest of the project. Cite it as: OpenIAP contributors, *Why the Commerce Protocol Draws Its Boundaries Where It Does*, version 1.0, 6 September 2026.

In short

A provider accepting purchase evidence does not authorize a particular user's current access. Evidence acceptance, account ownership, subscription lifecycle and current entitlement are four different questions, and an integration has to keep them apart across store APIs, provider implementations, transport bindings and events that arrive late.

OpenIAP Commerce Protocol is a vendor-neutral server-side contract that makes those boundaries explicit and turns a defined subset of them into checks a provider can be run against. GraphQL source defines the structure and machine-readable metadata, and the build generates the REST and GraphQL bindings, capability discovery and JSON Schemas from it. Normative prose defines the behavior, and the conformance vectors encode those rules by hand against the generated schemas. That hand step is where drift is most likely to enter.

The sections below cover the six boundaries, the two roles, the temporal meaning of an entitlement snapshot, event identity and recovery, and what the contract still leaves to you. What the current checks do and do not establish is recorded in the [evaluation record](#); the open questions and what would settle them are in the [research agenda](#).

1. Five symptoms

If you ship subscriptions, some of these have happened to you.

A paying customer is locked out because your verification provider was down and the code treated "no answer" as "not a valid purchase." A customer turns off auto-renewal on Tuesday and loses access on Tuesday, though they paid through the end of the month. A webhook arrives late, after the subscription it describes has already expired, and your handler opens the gate anyway because the payload says `active`. The same webhook arrives twice and the customer is credited twice. You move to a different verification provider and your event handling breaks, because nothing said what its events were supposed to mean.

None of these is a payment bug. Every one of them is a boundary that got crossed: a question was answered with the answer to a different question.




Figure 1. Four questions an integration must keep apart. Answering one does not answer the next.

The stores say as much themselves. Google's lifecycle guidance distinguishes cancellation before the end of a paid period from expiration, and describes continued access during grace [3]. Apple defines a notification's subscription status relative to the time the notification was signed [4]. What a message asserts and when that assertion applies are two different things.

The design question behind the specification follows from that: **how can a server contract preserve the meaning of purchase verification and current entitlements across provider and transport boundaries?** Everything below answers it for OpenIAP Commerce Protocol 1.0 [5], with IAPKit as the reference implementation.

The scope is principally auto-renewing subscriptions on Apple and Google Play, although the protocol's store identifier space is extensible. The contract does not prescribe every store transition, implement payment settlement, or amount to a complete commerce platform.

2. What this reasoning rests on

The boundaries below are not invented here. Two decades of results show integrations failing at the seams with the cryptography intact: logic flaws at the merchant-cashier boundary let shoppers pay nothing without breaking anything [12]; automatically rewriting apps so on-device checks returned success defeated in-app billing [11], and a later attack reached the same result by instrumenting the running app [1]; flawed third-party payment integrations were traced to SDK design, documentation and sample code [2], and a decade after the cashier-as-a-service results the same class of flaw was still being reported [13]. One result shaped this document's form more than the others: Chen et al. found payment-integration requirements a developer cannot enforce, because the parameters a check would need are not available to the party asked to perform it, and because guidance omits checks that are required [22]. The conclusion we draw from that is ours, not theirs: an obligation written only as advice is one an implementation cannot be failed against, so this contract states what can be checked.

The event rules come from the same place. Repeated messages need application-level handling and often remembered state, and an operation that is naturally idempotent is a different thing from one made idempotent [6], [14]. Idempotence alone still does not resolve missing events, stale snapshots or ambiguous ordering, which is why the contract addresses those separately.

Three of the design choices here have a source, though the choices remain ours. Deriving tests from a model rather than from an implementation is an established method [8], and generating stateful requests from a service's own specification finds real defects [15] — in that work, through server errors. The vectors here carry expected outcomes drawn from the normative prose, so they check required behavior rather than only request validity or a server error. Those outcomes are reviewed apart from the implementations because independent implementations can share an error, which is a finding [9]; reviewing them separately is our response to it, not a remedy that paper establishes.

And an unavailable verifier and an incomplete read get explicit rules rather than being left to each implementer, because catastrophic failures concentrate in already-signalled errors that were mishandled — 92% against 25% for non-catastrophic ones, in five distributed data systems [10]. That last inference is ours; those are different systems.

Two comparisons place the work. Bishop et al. wrote a behavioral specification of TCP and Sockets and checked real implementations against it [21] — the same shape as this, except theirs was validated against implementations they had not written, while this one is authored alongside one of its own. Pact is the contract testing you will compare this to [23]; its contracts are consumer-driven, recorded from what one consumer needs, whereas these obligations do not change when a new consumer appears. Event interchange has a vendor-neutral envelope of its own [7], which this contract is not a profile of.

3. Problem model and trust boundaries

3.1 Actors and observations

The model contains four actors: a store, a commerce provider, a developer backend, and an application. The store is the source of purchase evidence and lifecycle facts. The provider verifies evidence and exposes normalized operations and events. The developer backend owns application-account authorization and consumes the provider's account-scoped results. The application initiates purchases and may request account-free verification.

The contract distinguishes a transaction, a subscription, an event, and an entitlement. A transaction records an economic occurrence; a subscription describes an arrangement; an event records a change or observation; an entitlement expresses access at an evaluation time. A verification verdict is a separate observation about submitted evidence. These are conceptual boundaries, not interchangeable names for one boolean.

For discussion, let a verification outcome be one of:

```
V(evidence) = Accepted | Rejected | OperationError(code)
```

This notation summarizes the existing contract; it does not add a wire type. `Accepted` and `Rejected` correspond to successful operation results with `isValid: true` and `isValid: false`. If the provider cannot obtain a verdict, the relevant operation error is `VERIFICATION_FAILED`. Other conditions, such as authentication failure or rate limiting, retain their own categories. No outcome in this expression independently proves which application account owns the purchase.

3.2 Fault and adversary assumptions

The application may be modified, and a credential shipped inside it may be copied. Requests can contain malformed or mismatched evidence. An upstream verification call can fail without producing a verdict. Store notifications and provider webhooks can be delayed, repeated, or unavailable. A downstream consumer may observe a previously correct snapshot after its deadline.

The design assumes that the provider and developer backend enforce their declared roles and that actual verification establishes the required trust in store evidence. Conformance tests can expose specific violations of these obligations; they cannot establish that a malicious provider tells the truth or that a compromised backend protects its users. Cryptographic primitive security, stolen server credentials, and exhaustive store-verifier analysis are outside what any of the current checks address.

3.3 Decision-preservation obligations

Table 1 summarizes the selected obligations from the specification [5]. The specification remains authoritative for their complete wording and exceptions.

Boundary	Required distinction	Consequence for an integration
Evidence to verdict	Rejection differs from inability to obtain a verdict.	An unavailable verifier cannot supply a negative purchase verdict.
Verdict to account	Verification does not establish a binding.	Account association requires the authenticated backend's decision.
Lifecycle to access	Renewal intent differs from current entitlement.	Disabling renewal does not itself close an unexpired paid gate.
Snapshot to delivery	Evaluation time differs from arrival time.	A delayed grant cannot open access at or after its known expiry.
Repeated delivery to effect	An event identity is stable within its emitter context.	Reprocessing that event must not repeat substantive effects.
Provider to consumer	Compatible semantics need not imply identical identifiers.	Portability comparisons normalize only differences the contract permits.

These distinctions prevent particular category errors. Whether developers apply them correctly, and how often existing systems violate them, require empirical study beyond the contract's definition.

4. Protocol design

4.1 Authored contract and generated artifacts

The protocol separates structural authoring from behavioral obligations. GraphQL source files [16] define types, operations, and contract metadata; normative prose, using the RFC 2119 keywords [19], defines their meaning. The build generates the assembled contract, JSON Schemas in draft 2020-12 [17], operation artifacts, REST and GraphQL bindings, OpenAPI 3.1 documentation [18], and lifecycle vectors. Checks reject drift between authored inputs and checked-in outputs.

These artifacts are distributable without an OpenIAP account or a central runtime service. The portable runner imports no IAPKit implementation and uses caller-supplied adapters and validation support. A provider can expose supported profiles and bindings through capability discovery. Thus adopting the contract does not require adopting IAPKit's storage, deployment model, or credential format.

Single-source generation reduces opportunities for artifacts to disagree, but also creates a shared-fault risk. A mistaken authored rule can propagate consistently into every generated artifact. Drift checks and behavioral validation therefore answer different questions.

4.2 Operations and authorization

The six operations are `verifyPurchase`, `subscriptionStatus`, `entitlements`, `bindPurchase`, `eraseUser`, and `providerCapabilities`. Providers declare their supported profiles and must answer honestly for what they declare. A caller can branch on those declarations. A caller that skips discovery is not punished for it: an operation from a profile the provider serves is not refused on that ground, though it can still fail for any ordinary reason, and an operation from an unserved profile returns `UNSUPPORTED_PROFILE` rather than a wrong answer.

`verifyPurchase` is account-free. It accepts store-discriminated evidence and returns the authoritative `isValid` gate; its advisory `state` field does not replace that gate. The operation cannot read, select, or mutate account state. `bindPurchase`, by contrast, requires the server role, associates evidence with the backend's user identity, and does not transfer an existing binding. Possession of evidence alone is deliberately insufficient authority to bind it.




Figure 2. What each role may call. A verification credential must never reach an account read or mutation, which is what stops a shipped application from walking arbitrary user identities.

The server-only status and entitlement operations return tokenless views. They exclude purchase tokens, signed receipts, store transaction identities, and provider-internal record identifiers. Unknown optional information is omitted, though GraphQL cannot express omitted-versus-null on a selected member, so there it arrives as `null` and a caller normalizes that to absent. If an operation cannot provide the complete answer required by the contract, it fails rather than presenting a partial read as complete.

Authorization for privileged operations precedes operation-input validation, including GraphQL variable coercion. A resolver-only check is insufficient because coercion can execute first. The specification allows earlier transport-shape failures, such as an unparseable document. This ordering makes the role boundary observable and testable without making it dependent on the implementation's internal call structure.

4.3 Temporal entitlement semantics

For an event carrying a subscription snapshot, let `s` be its state, `d` its optional expiry, and `t` its `processedAt`. The provider evaluates the following predicate, restated from specification §2.3:

```

A(s, d, t) =
  (s is Active or InGracePeriod)
  and (d is absent or t < d)
  
```

For grace, `d` is the grace deadline. At equality the gate is closed. When expiry is absent, the predicate follows the state but supplies no deadline or general freshness guarantee. The provider places this result in `subscription.active`. Consumers read that decision instead of reconstructing it from `state` while ignoring `active`.

A consumer applying a true snapshot must additionally respect a present deadline at arrival and during subsequent use. It can schedule expiry or obtain an authoritative status before the deadline. These duties do not allow a false snapshot to become true merely because its state label looks active. When no deadline is available, freshness requirements need an appropriate authoritative status source.

Consider the illustrative trace in Figure 3. Times are synthetic and do not report a store experiment. Assume a bound subscription, no intervening revocation, and a known paid deadline at 10:00.




Figure 3. A snapshot that was true when it was evaluated does not become authorization when it arrives. A correct signature and an originally correct `active` value are insufficient to authorize access at a later time.

4.4 Events, identity, and recovery

The event path is store to provider to developer backend. The protocol defines bounded request/response operations for application-facing verification and server reads; its GraphQL binding has no Subscription root. Device notification delivery remains a developer-backend concern.

Provider delivery is duplicate-capable and unordered. An event can be accepted zero times if delivery permanently fails or its retry budget is exhausted. The emitter preserves event identity and body across retries, while transport signing, an HMAC-SHA256 [20] over the attempt's timestamp and the exact body bytes, uses the current attempt's timestamp. Consumers deduplicate within the emitter's identity context and handle substantive effects idempotently. Retries do not supply an exactly-once guarantee.

When a consumer can correlate a stable purchase, an older snapshot cannot overwrite newer state. This does not make every effect of an older event obsolete. Nor does it create a universal ordering key: a user identity or product identity alone may be insufficient, and equal occurrence timestamps provide no tiebreaker. A consumer unable to establish current state uses the declared status or entitlement operations, or an emitter-specific authoritative source if those operations are unavailable.

First binding introduces another temporal boundary. Previously unbound changes are coalesced into the current gate; an expired historical grant is not replayed as a new entitlement. Actionable entitlement events require both a user and product identity. These rules connect purchase observation to account access without making account association implicit in verification.

4.5 Portability and evolution

The contract defines compatibility at the level of supported profiles, bindings, and observable results. Protocol versions and package-distribution versions are distinct. Some value spaces and objects allow compatible additions, while closed tokenless results and error envelopes restrict what a provider can return.

REST and GraphQL need not produce byte-identical documents. Their envelope rules and treatment of unknown input members differ, so cross-binding comparison must use the defined semantics. It must still detect a binding that silently drops required information or changes an operation failure into a successful answer.

It is worth separating what this project claims from what it does not. The individual distinctions are established practice: Google's own lifecycle guidance already separates a cancellation from the end of access [3], and this document does not claim to have introduced that one or any of the others.

Provider switching has a narrower meaning than automatic migration. Consumers may preserve their domain logic under compatible contracts while changing credentials and endpoints. Historical data, purchase correlation, and overlapping deliveries remain operational responsibilities. Emitter identifiers are not globally interchangeable, and this version does not guarantee deduplication across a provider cutover.

5. What this does not give you

The same person wrote the specification, the generators, the tests, the reference implementation and this document. Consistency among them proves they agree with each other, not that they are right. Independent review of the store mappings and the expected outcomes is what would reduce that risk, and it has not happened.

The checks are finite and selected. There is no proof over every possible event sequence, no evidence from real store verification, and no demonstration that two independently operated providers interoperate. The [evaluation record](#) states exactly what the current tests do and do not establish; the [research agenda](#) states what would settle the open questions.

Some work stays with you, and one gap is worth naming precisely. For a read, the contract does compose: SPEC §4.3 returns every product the user may access right now, deduplicated, together with the subscription records whose open gates produced them, so a second still-active record is accounted for. The gap is in attribution and in the event model. The gate is per user and product, so an entitlement event carries both, but a product change moves one subscription from one product to another and the lifecycle vectors model that as a single gate transitioning rather than the outgoing product's rights closing and the incoming product's opening. Even when a refund identifies a transaction, nothing says which grant loses authority while a later period or a second purchase is still valid. Family sharing separates the purchaser from the beneficiary, and a promotional grant has no purchase behind it at all; neither appears in the contract. Account merging is declared out of scope. A single subscription per user avoids only the part that comes from holding several at once; a product change or a refunded earlier renewal still raises the question. So the aggregate is given to you and the attribution is not: which grant a right came from, and what a later correction does to it, are rules you write.

The contract also does not define a complete store transition machine, universal purchase correlation, or historical-data migration. Where an expiry or an event is missing you need an authoritative read; no envelope makes an old observation current.

6. When this is worth adopting

Start by reading Table 1 against the code you already have. That costs an hour and needs no adoption at all. Find where your verification call fails and confirm the failure does not become a rejection. Find your cancellation handler and confirm it does not close access before the paid period ends. Write one test that delivers a valid entitlement snapshot after its own expiry and assert your gate stays closed. Passing all three does not mean the integration is right; it means these three failures did not appear in the cases you exercised. Account authority, duplicate effects and portability are untested by them, and Table 1 is where to look next.

If one fails, the smallest useful step is still not adoption. For the verification case it is SPEC §4.1's separation of a rejected verdict from an operation error, so an unreachable verifier stops producing a negative purchase verdict. For the other two it is the entitlement predicate in SPEC §2.3, which SPEC §4.3 answers at a stated read time: evaluate access from state and expiry, and stop reading a lifecycle label as an access decision. Both are changes inside your own code.

Adopt the contract itself when you verify on a server for more than one store, and want one decision to hold across stores whose evidence, notifications and lifecycle vocabularies do not match, with the rules written somewhere that does not belong to your provider. If you also ship several client frameworks, the separate client contract in this repository addresses that side; the two are complementary and neither requires the other.

The alternatives are worth naming honestly. A store's own server API is authoritative but store-specific by construction, so using them directly leaves the cross-store decision to you. A hosted entitlement service makes that decision for you and documents it well. The mature ones publish their cancellation, expiration and grace semantics, their duplicate handling, and their versioning and deprecation commitments, and you can hold them to all of it. What that documentation describes is how to use one service. Its semantics are that service's own commitment rather than an obligation any other provider has taken on. This contract is the third option: the same decisions written as a specification any provider can implement, with one set of checks that runs against all of them. It is younger and less proven than either alternative, and it does not run anything for you. A relevant analogue in shape is TM Forum's Product Inventory Management API [24], a standard carrying a conformance profile and a test kit, though what it standardizes is product inventory and its lifecycle notifications for telecommunications rather than store purchase evidence and the entitlement decisions that follow from it.

Do not adopt it expecting a hosted service; that is IAPKit's job, and IAPKit is one implementation of this contract rather than the contract itself. Do not adopt it expecting the checks to prove your provider correct. They test a defined subset of obligations, which is a floor and not a guarantee.

The contract, the generated artifacts and the portable checks are published under the same terms as the rest of the project, so a second implementation can be built and held to them without asking anyone.

References

Scholarly metadata and reading status are maintained in [bibliography.md](#); the entries below reproduce what this document cites so that it can be read on its own.

[1] Collin Mulliner, William Robertson, and Engin Kirda. 2014. *VirtualSwindle: An Automated Attack Against In-App Billing on Android*. ACM AsiaCCS. DOI: 10.1145/2590296.2590335. [Publisher DOI](#).

[2] Wenbo Yang, Yuanyuan Zhang, Juanru Li, Hui Liu, Qing Wang, Yueheng Zhang, and Dawu Gu. 2017. *Show Me the Money! Finding Flawed Implementations of Third-party In-app Payment in Android Apps*. NDSS. [Paper](#).

[3] Google. *Subscription lifecycle*. Google Play Billing documentation. Vendor documentation; accessed 5 September 2026. [Lifecycle guidance](#).

[4] Apple. *status*. App Store Server Notifications documentation. Vendor documentation; accessed 5 September 2026 through its Markdown representation. [Status definition](#).

[5] OpenIAP contributors. *OpenIAP Commerce Protocol Specification 1.0*. Evaluated project artifact at commit `100557091077ddcd8e1f22128f80bd2b9747912`, accessed 5 September 2026. [Pinned specification](#). This is the subject of the study, not independent evidence of its effectiveness.

[6] Pat Helland. 2007. *Life beyond Distributed Transactions: an Apostate's Opinion*. CIDR. Position paper. [Paper](#).

[7] CloudEvents authors. *CloudEvents, Version 1.0.2*. Technical specification; accessed 5 September 2026. [Pinned specification](#).

- [8] Mark Utting, Alexander Pretschner, and Bruno Legeard. 2012. *A Taxonomy of Model-Based Testing Approaches*. *Software Testing, Verification and Reliability* 22(5), 297-312. DOI: 10.1002/stvr.456. [University record](#).
- [9] Jie Ma, Ningyu He, Jinwen Xi, Mingzhe Xing, Liangxin Liu, Jiushenzi Luo, Xiaopeng Fu, Chiachih Wu, Haoyu Wang, Ying Gao, and Yinliang Yue. 2026. *When Specifications Meet Reality: Uncovering API Inconsistencies in Ethereum Infrastructure*. *Proceedings of the ACM on Programming Languages* 10, OOPSLA1, Article 111. DOI: 10.1145/3798219. [Open version](#).
- [10] Ding Yuan, Yu Luo, Xin Zhuang, Guilherme Renna Rodrigues, Xu Zhao, Yongle Zhang, Pranay U. Jain, and Michael Stumm. 2014. *Simple Testing Can Prevent Most Critical Failures: An Analysis of Production Failures in Distributed Data-Intensive Systems*. *OSDI*, 249-265. [Publisher record](#) and [paper](#).
- [11] Daniel Reynaud, Eui Chul Richard Shin, Thomas R. Magrino, Edward X. Wu, and Dawn Song. 2012. *FreeMarket: Shopping for free in Android applications*. *NDSS*. [Programme entry](#).
- [12] Rui Wang, Shuo Chen, XiaoFeng Wang, and Shaz Qadeer. 2011. *How to Shop for Free Online: Security Analysis of Cashier-as-a-Service Based Web Stores*. *IEEE Symposium on Security and Privacy*, 465-480. DOI: 10.1109/SP.2011.26. [Publisher DOI](#).
- [13] Shangcheng Shi, Xianbo Wang, and Wing Cheong Lau. 2021. *Breaking and Fixing Third-Party Payment Service for Mobile Apps*. *ACNS, LNCS 12727*, 3-26. DOI: 10.1007/978-3-030-78375-4_1. [Publisher DOI](#).
- [14] Pat Helland. 2012. *Idempotence Is Not a Medical Condition*. *ACM Queue* 10(4), 30-46. DOI: 10.1145/2181796.2187821. [Publisher DOI](#).
- [15] Vaggelis Atlidakis, Patrice Godefroid, and Marina Polishchuk. 2019. *RESTler: Stateful REST API Fuzzing*. *ICSE*, 748-758. DOI: 10.1109/ICSE.2019.00083. [Publisher DOI](#).
- References 16 to 20 are the technical specifications the contract is built on: its artifact formats, its requirement keywords, and the message authentication its webhook signatures use. Reference 23 identifies a testing tool the text compares against, and 24 a standard it is measured against. None of them is evidence about OpenIAP.
- [16] GraphQL Foundation. *GraphQL Specification, October 2021 Edition*. Technical specification; accessed 6 September 2026. The repository does not pin an edition. [Specification](#).
- [17] JSON Schema. *JSON Schema Specification, draft 2020-12*. Technical specification; accessed 6 September 2026. The generated schemas declare this dialect. [Specification](#).
- [18] OpenAPI Initiative. *OpenAPI Specification 3.1.0*. Technical specification; accessed 6 September 2026. [Specification](#).
- [19] Scott Bradner. 1997. *Key words for use in RFCs to Indicate Requirement Levels*. *RFC 2119*. DOI: 10.17487/RFC2119. [RFC](#).
- [20] Hugo Krawczyk, Mihir Bellare, and Ran Canetti. 1997. *HMAC: Keyed-Hashing for Message Authentication*. *RFC 2104*. DOI: 10.17487/RFC2104. [RFC](#).
- [21] Steve Bishop, Matthew Fairbairn, Michael Norrish, Peter Sewell, Michael Smith, and Keith Wansbrough. 2005. *Rigorous specification and conformance testing techniques for network protocols, as applied to TCP, UDP, and Sockets*. *SIGCOMM*, 265-276. DOI: 10.1145/1080091.1080123. [Publisher DOI](#).

- [22] Yi Chen, Luyi Xing, Yue Qin, Xiaojing Liao, XiaoFeng Wang, Kai Chen, and Wei Zou. 2019. *Devils in the Guidance: Predicting Logic Vulnerabilities in Payment Syndication Services through Automated Documentation Analysis*. USENIX Security, 747-764. [Publisher record and paper](#).
- [23] Pact Foundation. *How Pact works*. Tool documentation; accessed 6 September 2026. [Documentation](#).
- [24] TM Forum. *Product Inventory Management API (TMF637), v5.0*. Open API specification; accessed 6 September 2026. [Specification](#).